

CMake Successor Build Systems: A World Tour of Build Systems Towards Better C++ Builds

DAMIEN BUHL &
ANTONIO DI STEFANO

Reliable, Fast and Safe Builds



Tipi.**build**



EngFlow ▶▶



Damien

co-founder, tipi.build

@daminetreg / damien@tipi.build



Antonio

DevEx Engineer, EngFlow

@TheGrizzlyDev / antonio@engflow.com

A love ❤️ letter to CMake

Giving CMake Superpowers



Thank you CMake 
autotools, qmake and SCons almost
disappeared 

Safety

C++ cares about Safety

- **Keynote: All the Safeties**, Sean Parent - C++ on Sea 2023 :
youtu.be/BaUv9sgLCP
- **C++ and Safety**, Timur Doumler - C++ on Sea 2023 :
youtu.be/imtpoc9jtOE
- **Safety and Security: The Future of C++**, JF Bastien - CppNow 2023 :
youtu.be/Gh79wcGJdTg

Are we there yet ?

“Are we there yet?”

- We have come a long way
 - From “classic C”
 - From “C with Classes”
 - From C++11
- We can write type-and-resource safe C++
 - Use contemporary C++
 - Avoid C-style and 1980s style C++
- C++ Core Guidelines
 - Directions, rules, and some enforcement
 - But not standardized
 - Uneven enforcement across implementations
- We need to standardize “Profiles”
 - And get implementations deployed
 - Finally reach type-and-resource safe C++!

Stroustrup - C++ safety - CppCon - October 2023



Are we there yet ?

Are We "Thing" Yet

C++

Are We Web Yet



Are We Game Yet



Are We Async Yet



Are We GUI Yet



Are we stack-efficient yet?



Are we safe yet?



Are we safe yet ?

We could !

It's more configuration work 

Are we safe yet ?

We should !



EUROPEAN
COMMISSION

Brussels, 15.9.2022
COM(2022) 454 final
2022/0272 (COD)

Proposal for a

REGULATION OF THE EUROPEAN PARLIAMENT AND OF THE COUNCIL
on horizontal cybersecurity requirements for products with digital elements and
amending Regulation (EU) 2019/1020

(Text with EEA relevance)

{SEC(2022) 321 final} - {SWD(2022) 282 final} - {SWD(2022) 283 final}

EN

EN

National Security Agency | Cybersecurity Information Sheet

Software Memory Safety

Executive summary

Modern society relies heavily on software-based automation, implicitly trusting developers to write software that operates in the expected way and cannot be compromised for malicious purposes. While developers often perform rigorous testing to prepare the logic in software for surprising conditions, exploitable software vulnerabilities are still frequently based on memory issues. Examples include overflowing a memory buffer and leveraging issues with how software allocates and deallocates memory. Microsoft® revealed at a conference in 2019 that from 2006 to 2018 70 percent of their vulnerabilities were due to memory safety issues. [1] Google® also found a similar percentage of memory safety vulnerabilities over several years in Chrome®. [2] Malicious cyber actors can exploit these vulnerabilities for remote code execution or other adverse effects, which can often compromise a device and be the first step in large-scale network intrusions.

Commonly used languages, such as C and C++, provide a lot of freedom and flexibility in memory management while relying heavily on the programmer to perform the needed checks on memory references. Simple mistakes can lead to exploitable memory-based vulnerabilities. Software analysis tools can detect many instances of memory management issues and operating environment options can also provide some protection, but inherent protections offered by memory safe software languages can prevent or mitigate most memory management issues. NSA recommends using a memory safe language when possible. While the use of added protections to non-memory safe languages and the use of memory safe languages do not provide absolute protection against exploitable memory issues, they do provide considerable protection. Therefore, the overarching software community across the private sector, academia, and the U.S. Government have begun initiatives to drive the culture of software development towards utilizing memory safe languages. [3] [4] [5]

U000219086.22 | PP-22-1723 | NOV 2022 Ver. 1.0

MARCH 2023

NATIONAL CYBERSECURITY STRATEGY

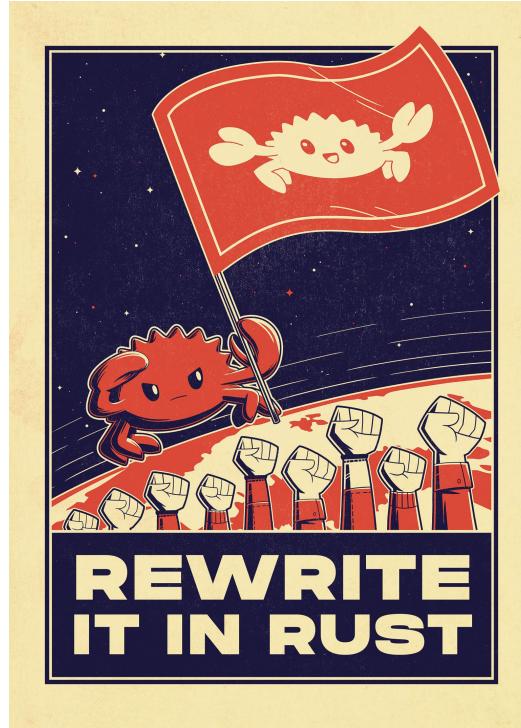
What is Safety ?

- Type Safety
- Memory Safety
 - Lifetime safety
 - Bounds safety
 - Initialization safety
- Object access safety
- Thread Safety
- Arithmetic safety
- Definition safety

Software Supply Chain Safety ?

- Reproducibility
- Traceability
- Release delivery time

YOU DON'T NEED TO



Achieve safety in C++ today

Achieve safety in C++ today
with the
CppCoreGuidelines



Achieving safety in C++ today
with **CMake**, **Clang**
Tidy and the
CppCoreGuidelines

```
set(CMAKE_CXX_CLANG_TIDY
    clang-tidy -checks=-*,cppcoreguidelines-*)
```

Clang Tidy Checks

Achieving safety in C++ today
with CMake

Achieving safety in C++ today

Let's tackle first

- Memory Safety
- Arithmetic Safety

Sanitizers for C++ builds with CMake

github.com/tipi-build/sanitizers-cmake

 **sanitizers-cmake** Public

forked from [arsenm/sanitizers-cmake](#)

 master  1 branch  0 tags

[Go to file](#) [Add file](#) [Code](#)

 README.md 

sanitizers-cmake

[open issues 10](#) [license MIT](#)

CMake module to enable sanitizers for binary targets.

Sanitizers for C++ builds

- AddressSanitizer (ASan)
- UndefinedBehaviorSanitizer (UBSSan)
- ThreadSanitizer (TSan)
- LeakSanitizer (LSan)
- MemorySanitizer (MSan)

Using LLVM Sanitizers in your CMake

```
project(safepp)

find_package(Sanitizers MODULE REQUIRED)

add_executable(use-after-free test/use-after-free.cpp)

# enable sanitizers
add_sanitizers(use-after-free)
```

So we done ?

**ASan, TSan, UBSan
are incompatible
with each others X**

Multiple builds and test runs

Separate CMake toolchain files

sanitize-ub.cmake :

```
set(SANITIZE_UNDEFINED TRUE CACHE BOOL "" FORCE)
```

sanitize-address.cmake :

```
set(SANITIZE_THREAD TRUE CACHE BOOL "" FORCE)
```

Multiple builds and test runs

Separate CMake toolchain files

sanitize-memory.cmake :

```
set(SANITIZE_MEMORY TRUE CACHE BOOL "" FORCE)
```

sanitize-thread.cmake :

```
set(SANITIZE_THREAD TRUE CACHE BOOL "" FORCE)
```

We want to run these Sanitizers

on each key strokes

**That requires a
decent amount of
CPU power**

Let's look at other build systems

Learning from other
build-systems

Parallelism

How many independent targets can a build system run concurrently at most?

Parallelism

Level 1 - Gradle

```
.  
├── app  
│   ...  
│   └── build.gradle  
└── lib  
    ...  
    └── build.gradle  
        settings.gradle
```

Parallelism

Level 1 - Gradle

Sub-project level parallelism. It is very limited as projects have very low granularity and are limited in numbers. It is disabled by default

Parallelism

Level 2 - Make, FASTBuild

Targets are more granular and the size of the job pool can be specified by a flag (-j/--jobs)

Parallelism

Level 3 – Bazel

Same as the level before, but resource-intensive targets can reduce the level of parallelism on-demand

Parallelism

Level 4 - Ninja

You can define separate job-pools each of which with a different size (eg: link=1 and cc=16)

Will it CMake?

Job pools can be used when compiling...

```
1 set_property(GLOBAL PROPERTY JOB_POOLS
2           compile=16
3           link=1
4           codegen=16)
5
6 set_property(TARGET atarget
7             PROPERTY JOB_POOL_COMPILE compile)
8
9 set_property(TARGET atarget
10            PROPERTY JOB_POOL_LINK link)
11
12 add_custom_target(protocgen
13                   COMMAND protoc --cpp_out=./out server.proto
14                   JOB_POOL codegen
15                   COMMENT "Generating server code from proto file"
```

Will it CMake?

... and linking ...

```
1 set_property(GLOBAL PROPERTY JOB_POOLS
2     compile=16
3     link=1
4     codegen=16)
5
6 set_property(TARGET atarget
7     PROPERTY JOB_POOL_COMPILE compile)
8
9 set_property(TARGET atarget
10    PROPERTY JOB_POOL_LINK link)
11
12 add_custom_target(protocgen
13     COMMAND protoc --cpp_out=./out server.proto
14     JOB_POOL codegen
15     COMMENT "Protobuf code generation"
```

Will it CMake?

... and pretty much anywhere else 😊

```
1 set_property(GLOBAL PROPERTY JOB_POOLS
2     compile=16
3     link=1
4     codegen=4)
5
6 set_property(TARGET atarget
7     PROPERTY JOB_POOL_COMPILE compile)
8
9 set_property(TARGET atarget
10    PROPERTY JOB_POOL_LINK link)
11
12 add_custom_target(protocgen
13     COMMAND protoc --cpp_out=./out server.proto
14     JOB_POOL codegen
15     . . .
```

Reproducibility

Given the same inputs and configuration, a target should yield identical outputs

Reproducibility

Level 1 - Make, Ninja, FASTBuild

Provide no real facility to ensure reproducibility
and, thus, hermeticity is something that the
developers have to take care of

Reproducibility

Level 2 – Meson

Allows you to query and introspect tools during a build and configure them to get reproducible results as well as define toolchains

Reproducibility

Level 3 – Gradle

Predefines robust and easy to configure toolchains out of the box. This way builds do not normally rely on any system-wide installed tools/libraries that can change between different workstations

Reproducibility

Level 4 - Bazel

Toolchains are treated as normal inputs. This trivially allows defining custom toolchains whilst keeping the same level of reproducibility

Will it CMake?

Actually, yes! But we have a few problems to fix



Will it CMake?

Problem 1: file ordering matters

Building `a.cpp` and `b.cpp` can yield different results from building `b.cpp` and `a.cpp`!

Will it CMake?

Problem 1: file ordering matters

Luckily CMake forces you to write input sources
manually, so ordering is guaranteed!

Will it CMake?

Problem 2: `__DATE__` and `__TIME__` macros

Will it CMake?

Problem 2: `__DATE__` and `__TIME__` macros

MSVC

```
add_link_options("/Bprepro")
```

Will it CMake?

Problem 2: `__DATE__` and `__TIME__` macros
GCC (via an environment variable)

```
SOURCE_DATE_EPOCH=1621012303
```

Will it CMake?

Problem 2: `__DATE__` and `__TIME__` macros

GCC + Clang

```
# Rewrite __DATE__
add_definitions(-D__DATE__="May 14 2021")
add_definitions(-D__TIME__="17:11:43")
add_compile_options(-Wno-builtin-macro-redefined)
```

Will it CMake?

Problem 3: Randomness when using `--flto` with
GCC

LTO on GCC will produce random symbols

Will it CMake?

Problem 3: Randomness when using `--flto` with GCC

```
set(APP_SOURCES
    app.cpp
    lib.cpp)
foreach(_file ${APP_SOURCES})
    file(SHA1 ${_file} shalsum)
    string(SUBSTRING ${shalsum} 0 8 shalsum)
    set_property(SOURCE ${_file}
        APPEND_STRING PROPERTY
        COMPILE_FLAGS
        "-frandom-seed=0x${shalsum}" )
endforeach()
```

Will it CMake?

Problem 4: `__FILE__` macro

The `__FILE__` macro is very likely to change between different workstations and thus produce different results

___FILE___ macro and path prefix

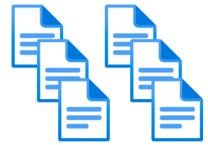
- Just build in the same environment
- At the same path
- But how ?

Let's merge Git and Ninja !

Because we can use Git :

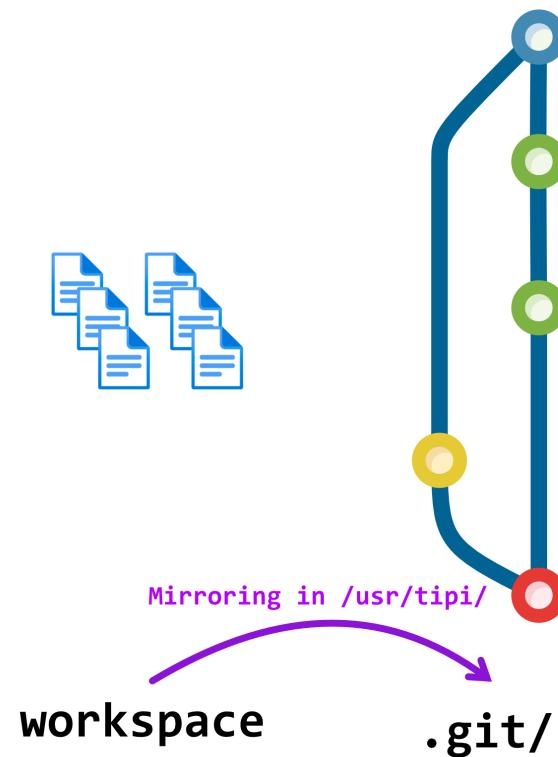
- Content Addressable Filesystem 
- to clone sources and build trees in invariant paths 
- to **track changes** 

Stores code in an invariant path



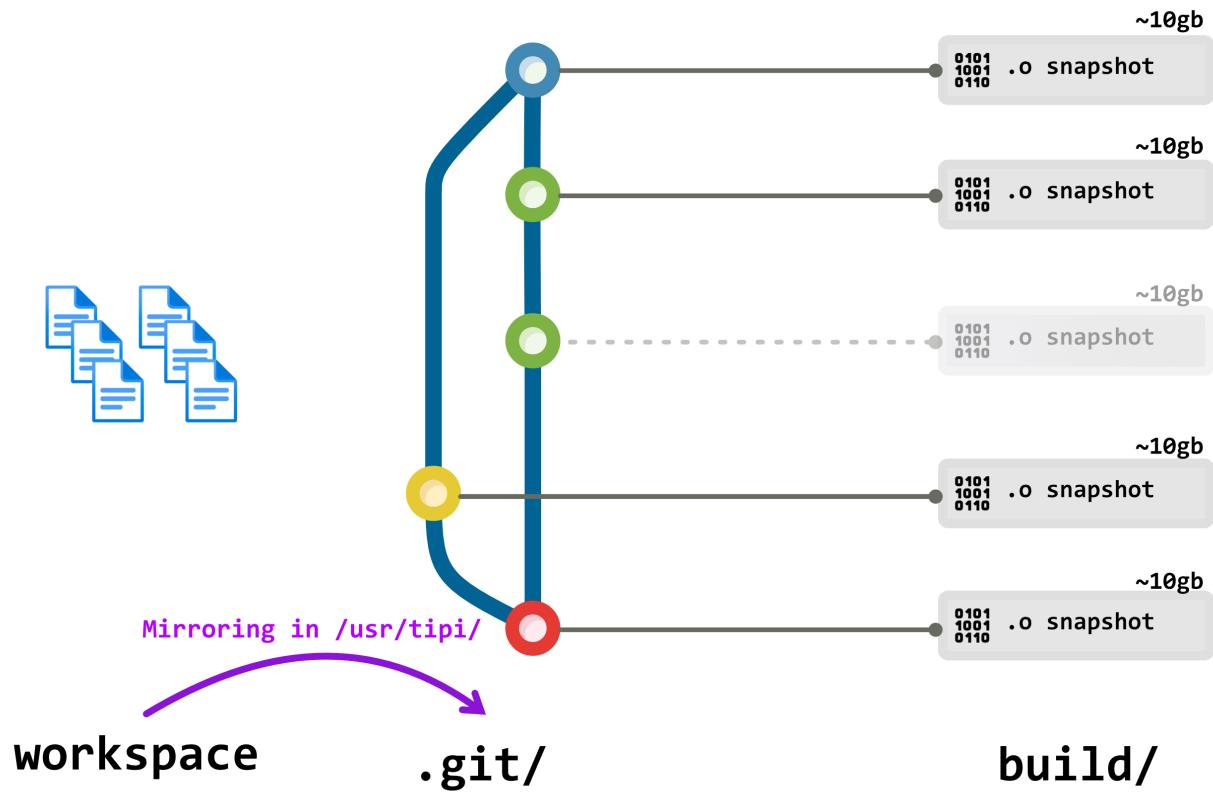
workspace

git based sources mirroring



Cache id based on the repository origin

commit-id addressable builds



Git + Ninja + cache packs

--FILE-- problem solved 🎉

Caching

How many targets can a build system avoid rebuilding?

Caching

Level 1 - Gradle

Once again the low level of granularity means that even small changes will rebuild very large targets

Caching

Level 2 - Make, Ninja

Check timestamps on the individual input files and rebuild targets if any change. Caching is completely file-system based and uses the mtime of a file to detect if it has changed since the last build

Caching

Level 3 – FASTBuild

Supports the same type of file-system caching as the previous level and adds support for distributed caching

Caching

Level 4 - Bazel

Supports file-system and distributed caching as with the previous level, but based on the digest of the input files' content rather than mtime. Bazel also has an in-memory cache for large build graphs

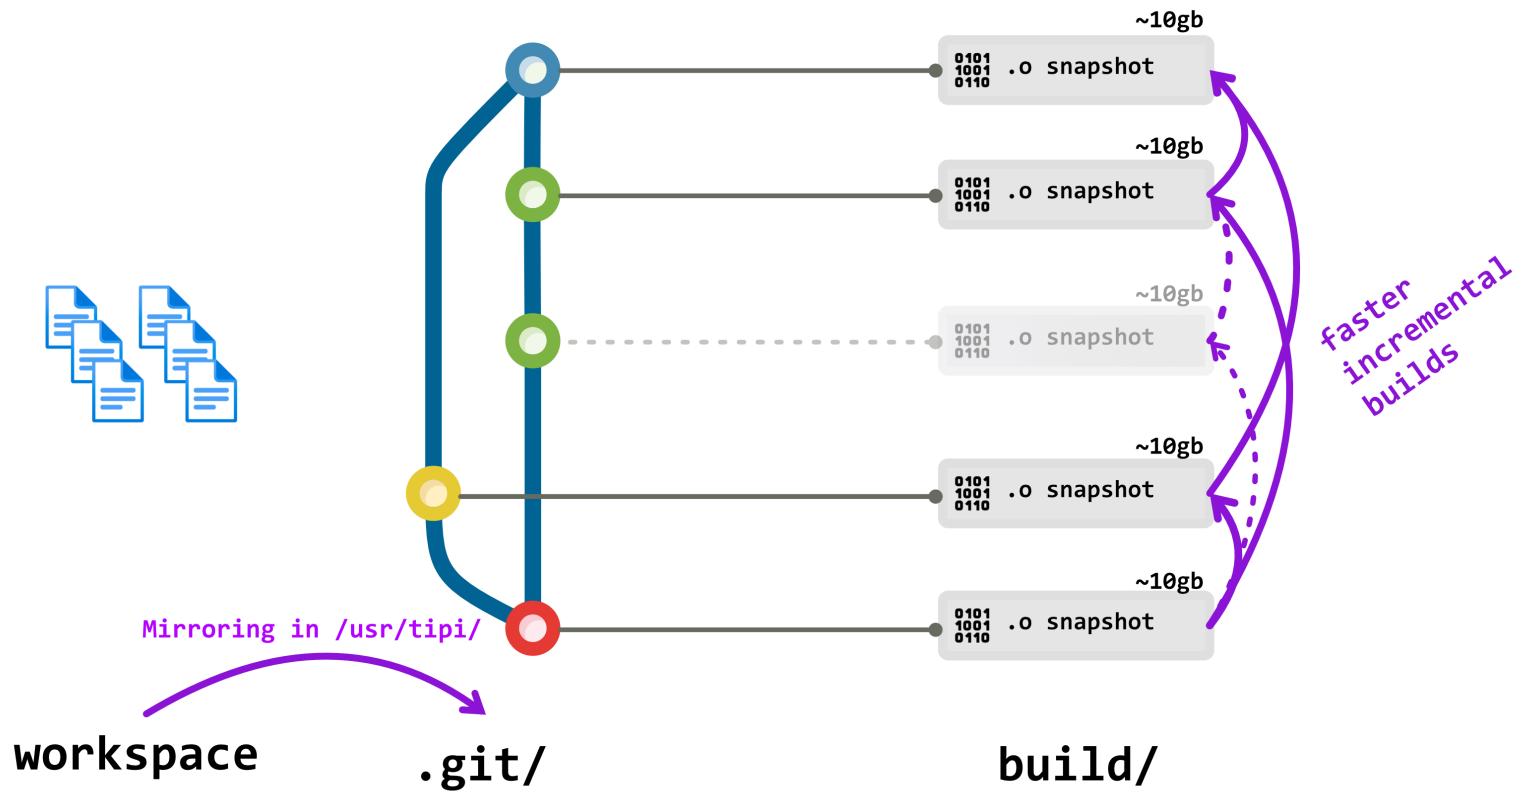
Will it CMake?

By using Make or Ninja we get file-system
caching out-of-the-box

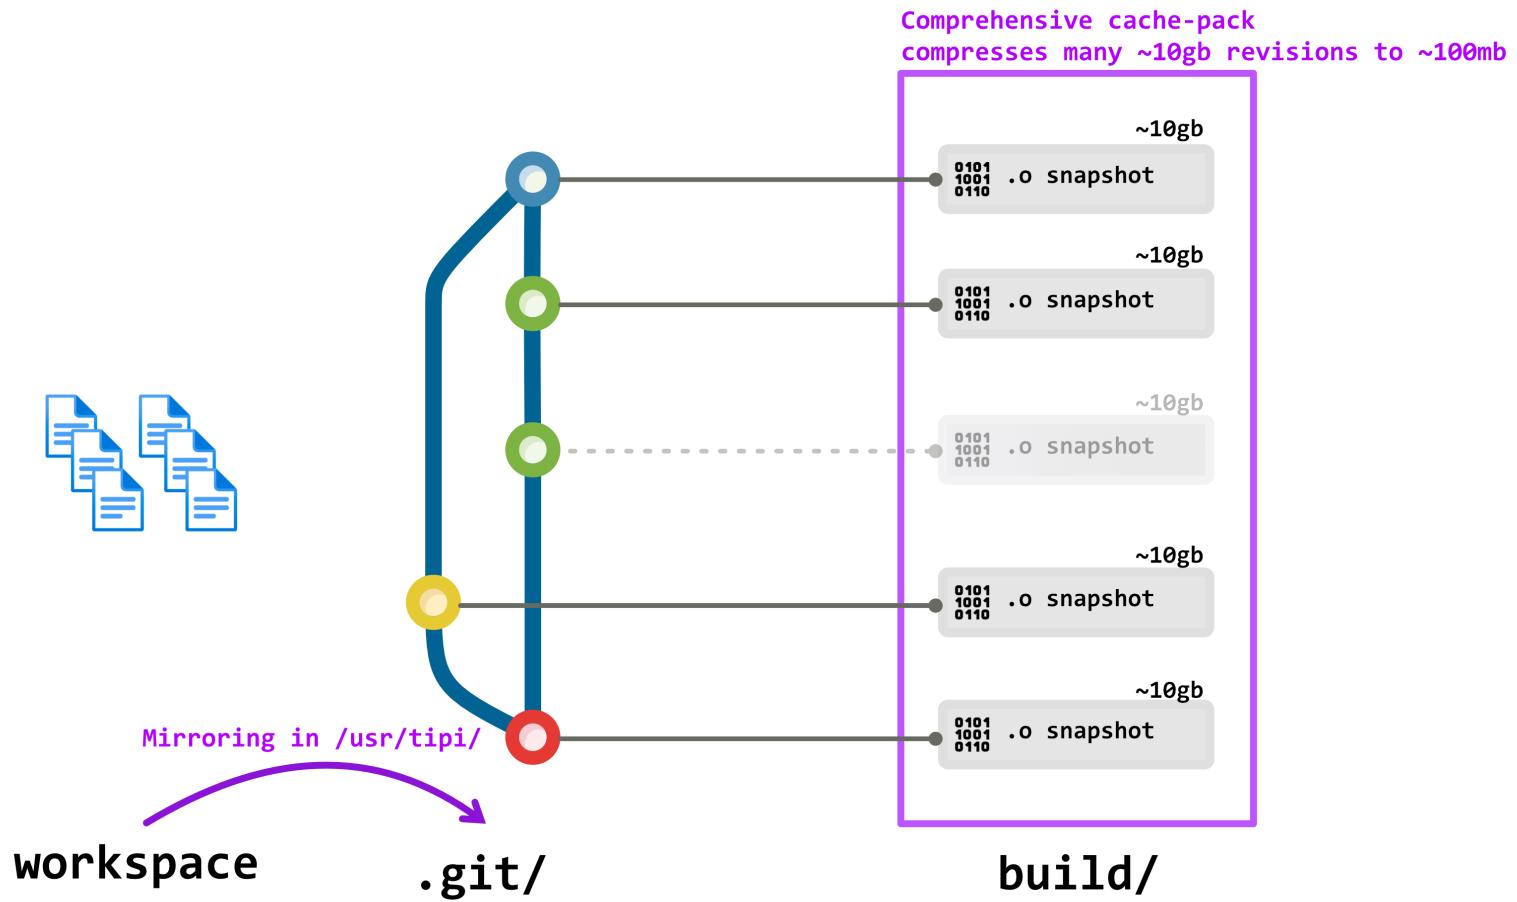
Will it CMake?

However, we can replace the standard CMake client with tipi. Engflow and Tipi partnered to bring remote execution and caching to CMake!

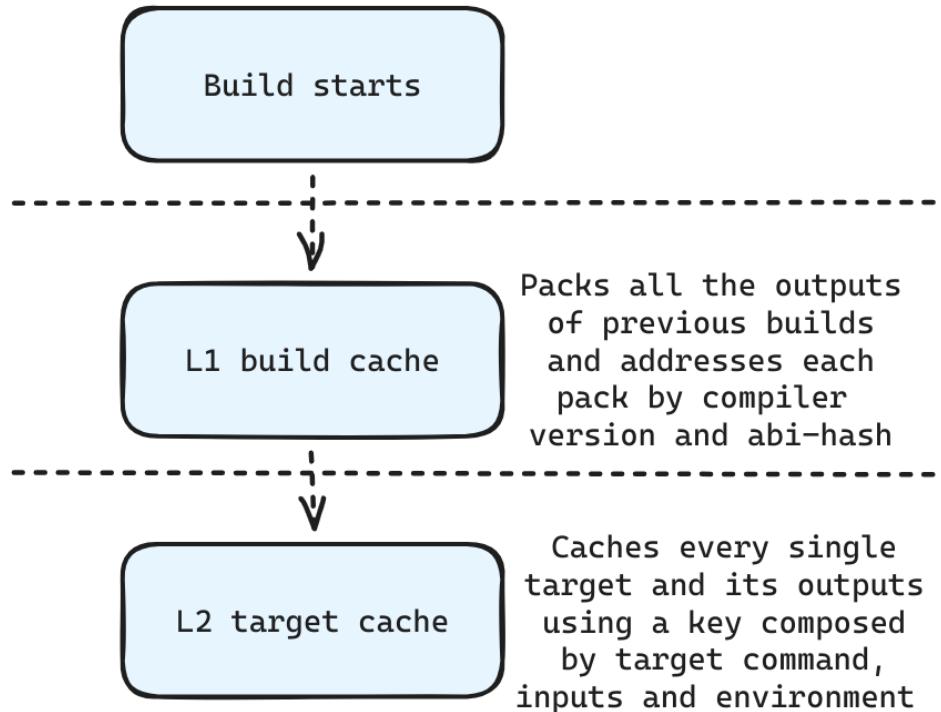
L1 Cache



L1 Cache



L2 Cache



Rationale

- L1 cache
 - permits fast build / install tree clean restore
- L2 cache
 - less cores required for parallel CI jobs

Dependency management

How a build system handles
3rd-party code

Dependency management

Level 1 - Make, Ninja, FASTBuild

No support at all, the developer gets to decide
how to handle dependencies

Dependency management

Level 2 - Gradle

Dependency management focuses on handling jars but is very crude everywhere else

Dependency management

Level 3 – Meson, Bazel

Both provide somewhat extensible mechanisms to resolve dependencies for multiple languages and both provide centralised repositories of dependencies tweaked to better work with the respective build systems

Will it CMake?

But package
management in C++
is 🙄

Actually we have a good solution

- CMake is the de-facto standard
- It has FetchContent ❤️

FetchContent

```
Include(FetchContent)
FetchContent_Declare(
    Boost
    GIT_REPOSITORY https://github.com/boostorg/boost.git
    GIT_TAG        boost-1.80.0
)
FetchContent_MakeAvailable(Boost)
find_package(boost_filesystem CONFIG REQUIRED)

target_link_libraries(app Boost::filesystem)
```

FetchContent is kinda slow



FetchContent is kinda slow



- ⚡ No problem : we can optimize it
- Let's use CMake

SET_DEPENDENCY_PROVIDER

Intercept FetchContent calls

```
macro(tipi_provide_dependency method package_name)
    set(oneValueArgs
        GIT_REPOSITORY
        GIT_TAG
    )
    cmake_parse_arguments(
        ARG "${options}" "${oneValueArgs}"
        "${multiValueArgs}" ${ARGN} )
    # Do something with ${ARG_GIT_REPOSITORY}
    FetchContent_SetPopulated( ${package_name} )
endmacro()
```

Install dependency provider

dependency_provider.cmake

```
cmake_language(  
    SET_DEPENDENCY_PROVIDER tipi_provide_dependency  
    SUPPORTED_METHODS  
    FETCHCONTENT_MAKEAVAILABLE_SERIAL  
)
```

```
list (APPEND CMAKE_PROJECT_TOP_LEVEL_INCLUDES  
    dependency_provider.cmake)
```

Let's install Boost

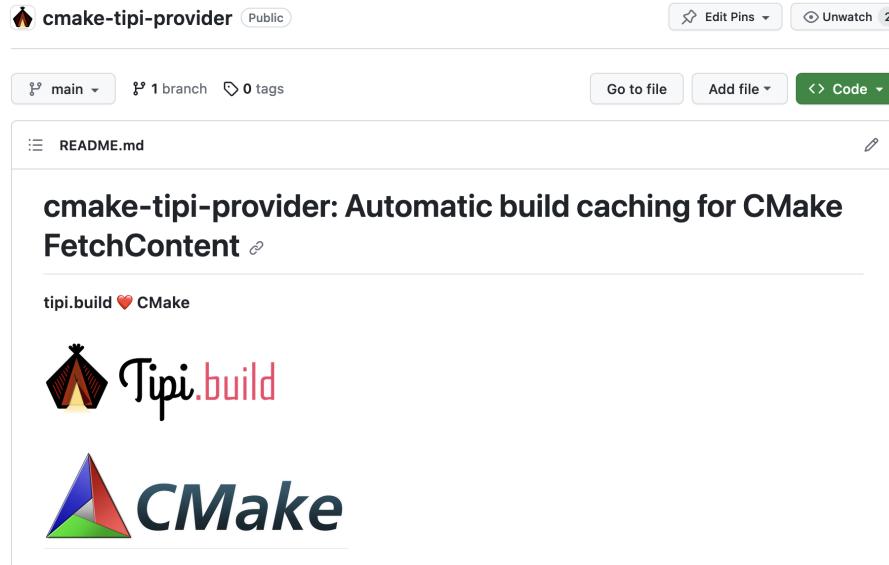
```
Include(FetchContent)
FetchContent_Declare(
    Boost
    GIT_REPOSITORY https://github.com/boostorg/boost.git
    GIT_TAG          boost-1.80.0
)
FetchContent_MakeAvailable(Boost)
```

Advantages

- No lock-in, pure CMake
- Rebuild from sources if needed
- Efficient and secure build caching

Automatic CMake Packages Caching

github.com/tipi-build/cmake-tipi-provider



**Build Provenance & Traceability
SBOMs**

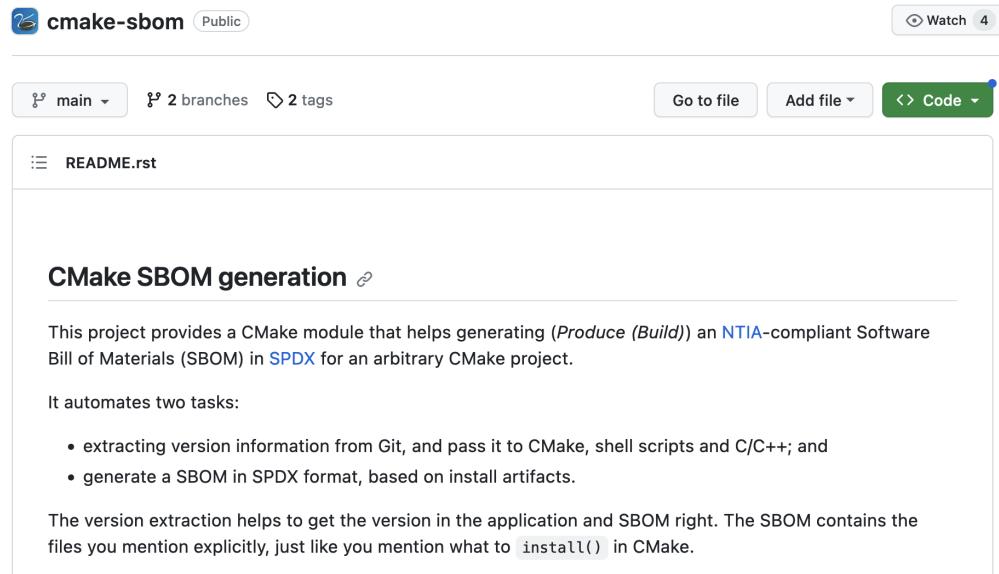
Build Provenance & Traceability

SBOMs

- Software Supply Chain Safety
- SPDX Format
- NTIA Compliance Validation

Manual Guided SBOM Generation

github.com/DEMCON/cmake-sbom



The screenshot shows the GitHub repository page for 'cmake-sbom'. At the top, there's a navigation bar with a logo, the repository name 'cmake-sbom' (marked as 'Public'), and a 'Watch' button with a count of 4. Below the header, there are buttons for 'main' (with a dropdown arrow), '2 branches', '2 tags', 'Go to file', 'Add file', and a green 'Code' button with a dropdown arrow. The main content area shows the 'README.rst' file. The title 'CMake SBOM generation' is bolded, followed by a link icon. A descriptive paragraph explains the project's purpose: 'This project provides a CMake module that helps generating (Produce (Build)) an NTIA-compliant Software Bill of Materials (SBOM) in SPDX for an arbitrary CMake project.' It then lists two tasks: extracting version information from Git and generating a SBOM in SPDX format based on install artifacts. A note at the bottom states: 'The version extraction helps to get the version in the application and SBOM right. The SBOM contains the files you mention explicitly, just like you mention what to `install()` in CMake.'

SBOMs generation

```
include(sbom)

sbom_generate(
    OUTPUT
        ${CMAKE_INSTALL_PREFIX}/sbom-${GIT_VERSION_PATH}.spdx
    LICENSE MIT
    SUPPLIER tipi
    SUPPLIER_URL https://tipi.build
)

reuse_spdx()

add_executable(app app.cpp)
```

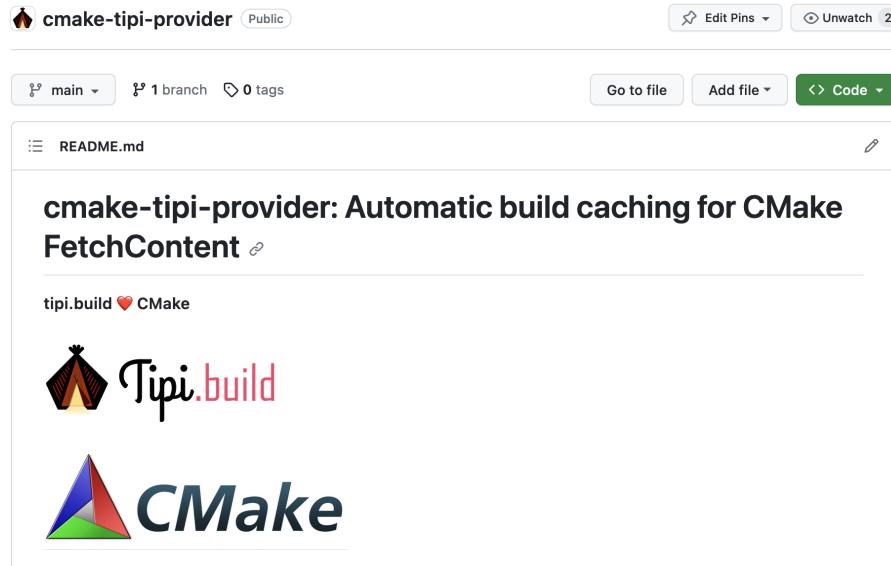
Automated SBOMs generation

```
install(
    TARGETS app
    EXPORT "${targets_export_name}"
    RUNTIME DESTINATION "bin"
)
sbom_add(TARGET app)

sbom_finalize()
```

Automatic SBOM Generation for your dependencies

github.com/tipi-build/cmake-tipi-provider



Build Cache Registry

View SBOMs, Dependencies, Cached builds

Distributed/remote builds

The ability of a build system to
cache and/or run targets on a
shared distributed system

Distributed/remote builds

The ability of a build system to
cache and/or run targets on a
shared distributed system

* also known as the thing that pays my salary :)

Distributed builds

Level 1 - Make, Ninja, Meson

No distributed caching or execution out of the box, everything runs locally

Distributed builds

Level 2 - Gradle

Has support for remote caching, but no remote execution of build targets. Once again the low level of granularity negatively impacts the potential of this feature

Distributed builds

Level 3 – FASTBuild

Supports both distributed execution and caching, however it is limited to just a list of well-known compilers when it comes to remote execution

Distributed builds

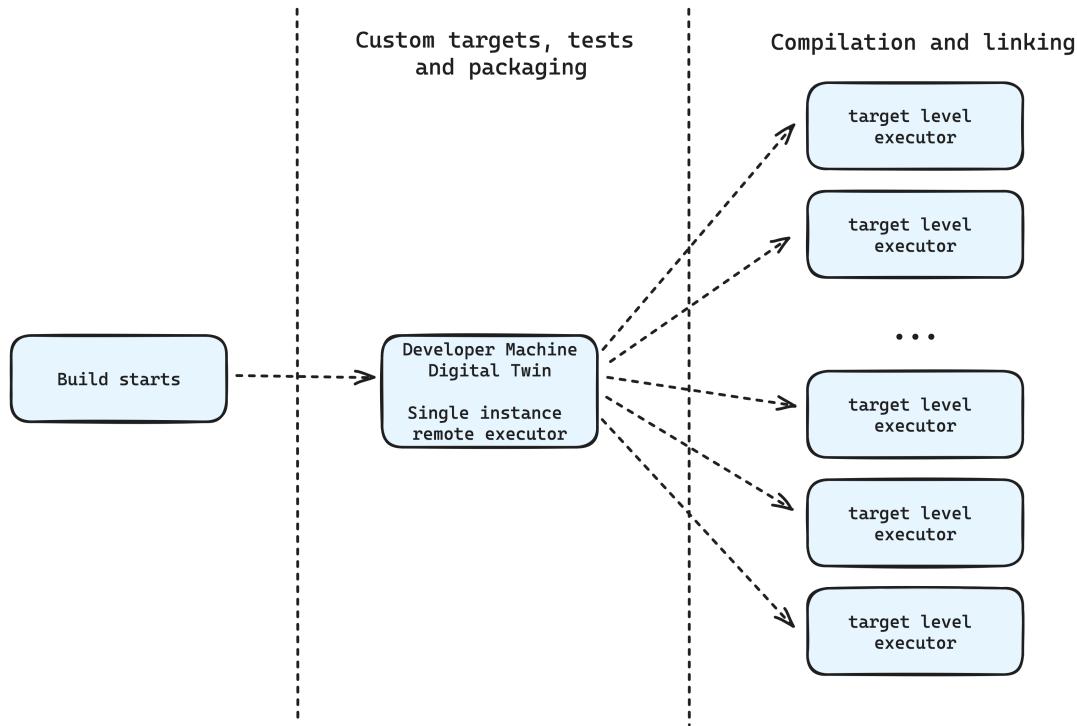
Level 4 - Bazel

Supports both distributed builds and caching for arbitrary tools, including compilers, linkers, tests and whatever you can think of

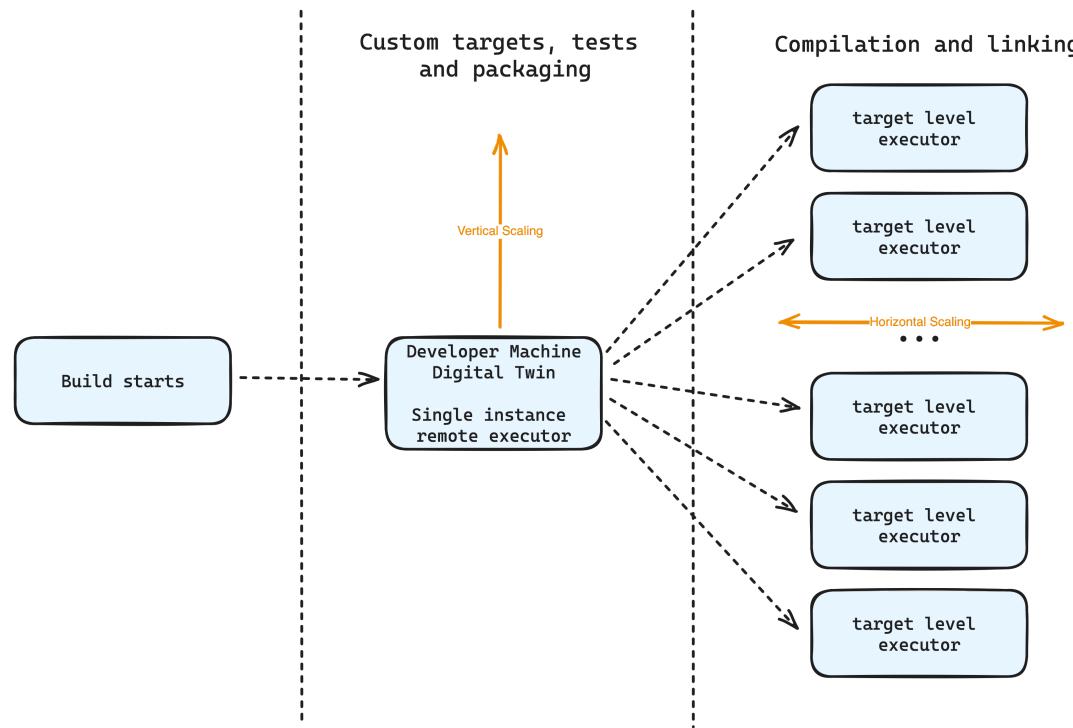
Will it CMake?

100%! We can combine Engflow's Remote Execution, which uses the same RE level protocol Bazel adopts, and tipi's single instance approach.

Will it CMake?



Will it CMake?



Let's check !

Thank you !

Request access to the BETA
<https://tipi.build/cmake-re>

<https://engflow.com/cmake-re>



